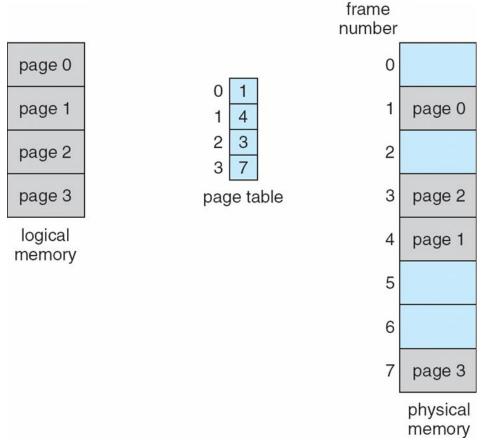
Paging Chap 18, 19, 20

 Memory management scheme that permits the physical address space of a process to be noncontiguous

Implemented by closely integrating the hardware and operating

system



Basic method

- Frames (page frames)
 - Breaking physical memory into fixed-size blocks
- Pages
 - Breaking logical memory into blocks of the same size
 - Page size
 - Typically power of 2
 - $-512B \sim 16MB$ (typically 4KB $\sim 8KB$)
- Logical address
 - v = (p, d)
 - p: page number
 - d: page offset (displacement)

page number	page offset
р	d
m -n	n

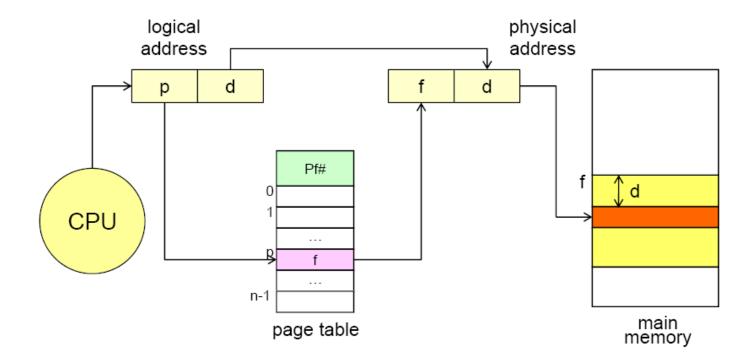
- logical address space 2^m
- page size 2ⁿ

Basic method

- Address mapping
 - Logical address → physical address
 - Hidden from the user and controlled by the operating system
 - Uses page map table (PMT)
 - A page table for each process
 - PMT resides in kernel space of main memory

Basic method

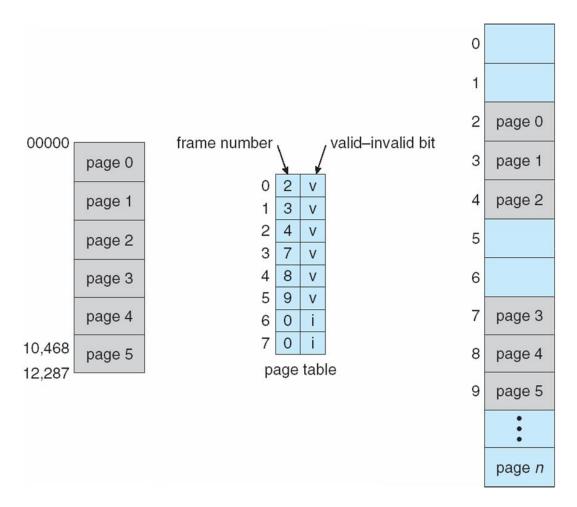
- Paging hardware and address mapping
- Page-table base register (PTBR) points to the page table
- Page-table length register (PTLR) indicates size of the page table



Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- Valid-invalid bit attached to each entry in the page table:
 - "valid" indicates that the associated page is in the process' logical address space, and is thus a legal page
 - "invalid" indicates that the page is not in the process' logical address space
 - Or use page-table length register (PTLR)
- Any violations result in a trap to the kernel

Valid (v) or Invalid (i) Bit In A Page Table

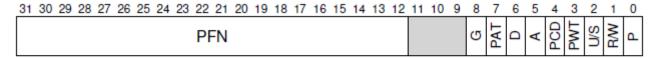


Page-table entry (PTE)

Valid bit

- indicate whether the particular translation is valid
- All the unused space will be marked invalid, and if the process tries to access such memory, it will generate a trap
- crucial for supporting a sparse address space
- remove the need to allocate physical frames for unallocated pages and thus save a great deal of memory.
- Protection bit (R/W, U/S)
- Present bit (P)
 - Indicates whether this page is in physical memory or on disk
- Dirty bit (D)
 - Indicates whether the page has been modified
- Reference bit (accessed bit) (A)
 - used to track whether a page has been accessed

No valid bit?



PWT, PCD, PAT, and G determine how hardware caching works

Page Table Location

Direct mapping

- Doubles the number of memory accesses during execution of the program - One for the page table and one for the data/instruction
- Performance degradation

TLB (Translation Look-aside Buffer)

- PMT: maintained in kernel space of the memory
- Subset of PMT entries: maintained in small TLB
 - Entries of the recently used pages
 - dedicated registers or cache memories

Memory Trace of Direct mapping

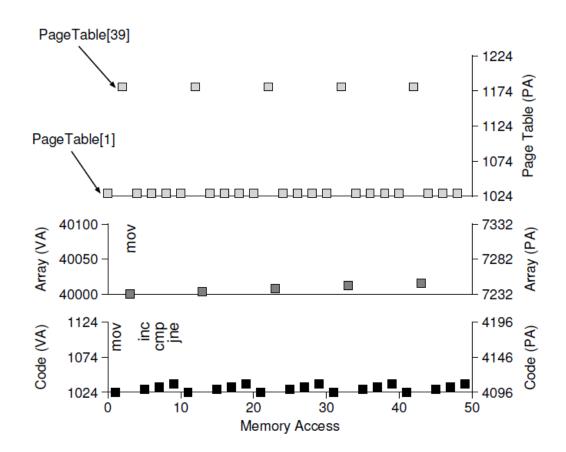


int array[1000]; ... for (i = 0; i < 1000; i++) array[i] = 0;



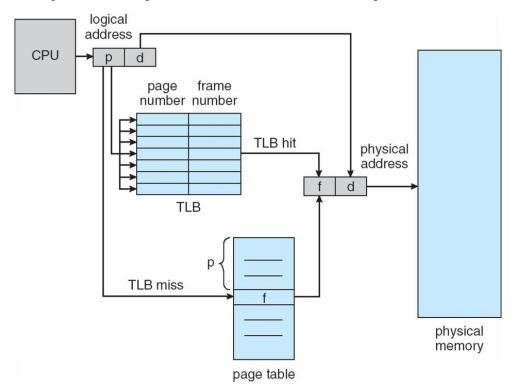
1024 movl \$0x0,(%edi,%eax,4) 1028 incl %eax 1032 cmpl \$0x03e8,%eax 1036 jne 0x1024





TLB (Translation Look-aside Buffer)

- Associative memory in MMU parallel search
- Address translation (p, d)
 - If p is in associative register, get frame # out
 - Otherwise get frame # from page table in memory
 - TLBs typically small (64 to 1,024 entries)



11

Effective Access Time

- TLB hit ratio $\alpha = 80\%$
- TLB search: 20 ns
- memory access: 100 ns
- EAT = $0.80 \times (20+100) + 0.20 \times (20+100+100) = 140$ ns
 - 40% slowdown in memory access time

Who Handles The TLB Miss?

```
VPN = (VirtualAddress & VPN_MASK) >> SHIFT
    (Success, TlbEntry) = TLB_Lookup(VPN)
    if (Success == True) // TLB Hit
        if (CanAccess(TlbEntry.ProtectBits) == True)
            Offset = VirtualAddress & OFFSET_MASK
            PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
            Register = AccessMemory(PhysAddr)
        else
            RaiseException (PROTECTION_FAULT)
    else
                          // TLB Miss
10
        PTEAddr = PTBR + (VPN * sizeof(PTE))
11
        PTE = AccessMemory(PTEAddr)
12
        if (PTE. Valid == False)
13
            RaiseException (SEGMENTATION_FAULT)
        else if (CanAccess(PTE.ProtectBits) == False)
            RaiseException(PROTECTION FAULT)
        else
            TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
            RetryInstruction()
19
```

software-managed TLB (RISC)

On a TLB miss, hardware simply raises an exception

RaiseException(TLB_MISS)

Need to replace an old one (LRU, Random)

hardware-managed TLB (CISC)

Hardware has to know exactly where the page tables are located in memory (via PTBR), as well as their exact format; on a miss. Hardware "walks" the page table, find the correct page-table entry and extract the desired translation, update the TLB with the translation, and retry the instruction

Software-managed TLB

- When returning from a TLB miss-handling trap, the hardware must resume execution at the instruction that caused the trap
 - This retry thus lets the instruction run again, this time resulting in a TLB hit.
- When running the TLB miss-handling code, the OS needs to be extra careful not to cause an infinite chain of TLB misses to occur.
 - Keep TLB miss handlers in physical memory (where they are unmapped and not subject to address translation),
 - Reserve some entries in the TLB for permanently-valid translations
- Flexible
 - OS can use any data structure it wants to implement the page table, without necessitating hardware change.
- Simple
 - hardware doesn't have to do much on a miss; it raises an exception, and the OS TLB miss handler does the rest.

Dongkun Shin, SKKU $oxed{1}^{2}$

TLB Valid Bit # Page Table Valid Bit

TLB entry also has several status bits

valid, protection, dirty, ...

Invalid in PTE

- The page has not been allocated by the process
- The access to an invalid page causes a trap to the OS

Invalid in TLB

- a TLB entry has not a valid translation within it
- When a system boots, a common initial state for each TLB entry is to be set to invalid, because no address translations are yet cached there.
- Context switch will set all TLB entries to invalid

TLB Issue: Context Switches

- TLB contains virtual-to-physical translations that are only valid for the currently running process
- When switching from one process to another, the hardware or OS (or both) must be careful to ensure that the about-to-be-run process does not accidentally use translations from some previously run process.
- Solutions
 - Flush the TLB on context switches, sets all valid bits to 0
 - incur TLB misses when a new process runs → high cost
 - Address space identifier (ASID) \approx (PID)
 - enable sharing of the TLB across context switches

		I	valid	L .		VPN	PFN	valid	prot	ASID
P1 →	10	100	1	rwx	-			1		
		_	0			_	_	0 1		
P2 →	10	170	1	rwx		10	170	1	rwx	2
	_	_	0	_		_	_	0	_	_

Shared Pages

Shared code

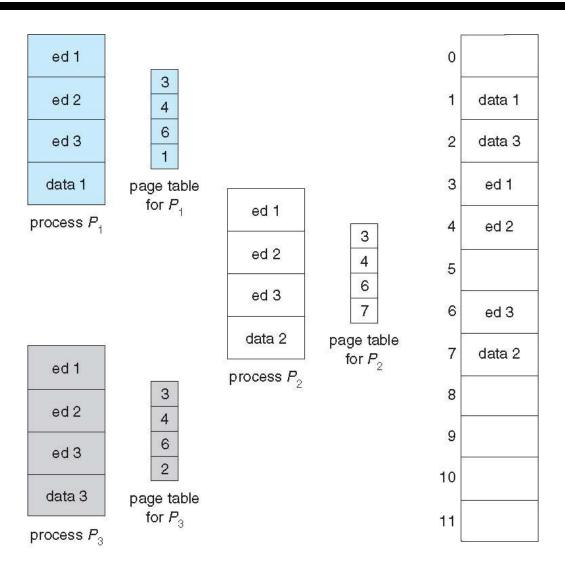
- One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of readwrite pages is allowed

Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space

VPN	PFN	valid	prot	ASID
10	101	1	r-x	1
_	_	0	_	_
50	101	1	r-x	2
_	_	0	_	_

Shared Pages Example



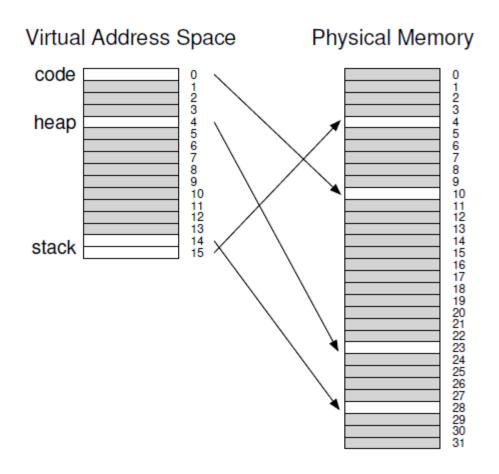
Paging: Smaller Tables

- Memory structures for paging can get huge using straightforward methods
 - Consider a 32-bit logical address space as on modern computers
 - Page size of 4 KB (2¹²)
 - Page table would have 1 million entries (2³² / 2¹²)
 - If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
 - That amount of memory used to cost a lot
 - Don't want to allocate that contiguously in main memory
- Bigger Page
- Hybrid Approach: Paging and Segment
- Multi-level Page Table (Hierarchical Paging)
- Hashed Page Table
- Inverted Page Table

Simple Solution: Bigger Pages

- 16KB pages (2¹⁴)
 - $-(2^{32}/2^{14})*4B=2^{20}B=1MB$
 - Large TLB coverage
- Internal fragmentation
- Most systems use relatively small page sizes
 - 4KB (x86) or 8KB (SPARCv9)
- Many architectures now support multiple page sizes
 - MIPS, SPARC, x86-64

Sparse Page Table



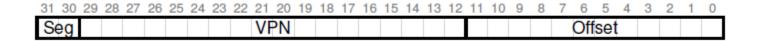
Most of the page table is unused, full of invalid entries

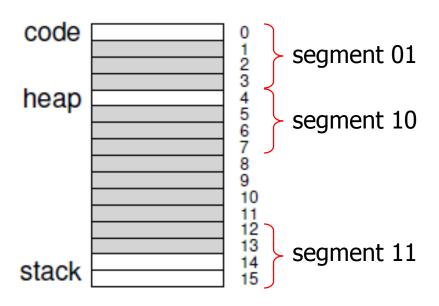
PFN	valid	prot	present	dirty
10	1	r-x	1	0
10		1-X	1	U
-	0		-	-
-	0	_	-	-
-	0		-	-
23	1	rw-	1	1
-	0		-	-
-	0		-	-
-	0		-	-
-	0		-	-
-	0		-	-
-	0		-	-
-	0		-	-
-	0		-	-
-	0		-	-
28	1	rw-	1	1
4	1	rw-	1	1

A 16KB Address Space With 1KB Pages

Page Table

Hybrid Approach: Paging and Segments



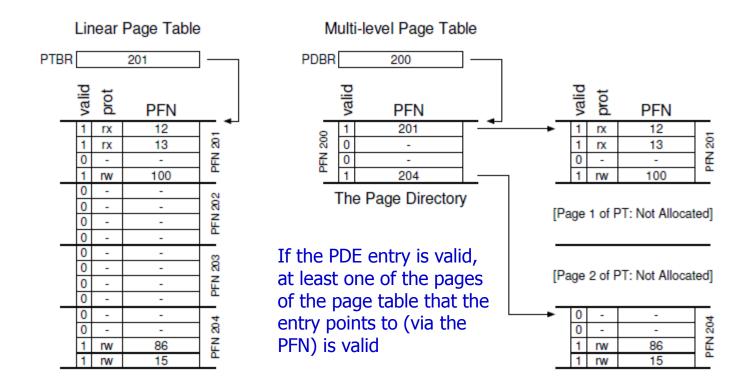


- The base register for each of these segments contains the physical address of a linear page table for that segment.
- Each process in the system now has three page tables associated with it.
- Page tables now can be of arbitrary size
- Unallocated pages between the stack and the heap no longer take up space in a page table
- Not flexible: if we have a large but sparsely-used heap,
 we can still end up with a lot of page table waste

External fragmentation

Multi-level (Hierarchical) Page Tables

- Break up the logical address space into multiple page tables
- Multi-level paging, in which the page table is also paged
- A simple technique is a two-level page table



Two-Level Paging Example

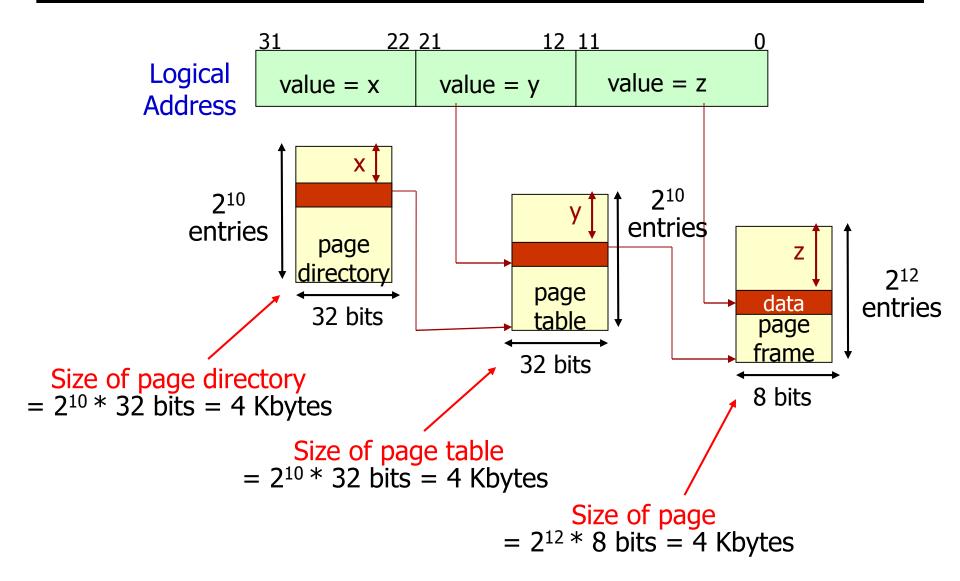
- A logical address (on 32-bit machine with 4K page size) is divided into:
 - a page number consisting of 20 bits
 - a page offset consisting of 12 bits (4KB page)
- Since the page table is paged, the page number is further divided into:
 - a 10-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:



• where p_1 is an index into the page directory, and p_2 is the displacement within the page of the page table

Dongkun Shin, SKKU 2^{2}

Address-Translation Scheme



Two-Level Page Table

Small (typical) system

- One page table might be enough
 - Page directory size + Page table size = 8 Kbytes of memory would suffice for virtual memory management
- How much *physical* memory could *this one page table* handle?
 - Number of page tables * Number of page table entries * Page size =

$$1 * 2^{10} * 2^{12}$$
 bytes = 4 Mbytes

Large system

- You might need the maximum number of page tables
 - Max number of page tables * Page table size =
 2¹⁰ directory entries * 2¹² bytes = 2²² bytes = 4 Mbytes of memory would be needed for virtual memory management
- How much *physical* memory could *these 2¹⁰ page tables* handle?
 - Number of page tables * Number of page table entries * Page size = 2¹⁰ * 2¹⁰ * 2¹² bytes = 4 Gbytes

Pros/Cons of multi-level page tables

- Only allocates page-table space in proportion to the amount of address space you are using
 - compact and supports sparse address spaces
- if carefully constructed, each portion of the page table fits neatly within a page, making it easier to manage memory
 - For a large page table (say 4MB), finding such a large chunk of unused contiguous free physical memory can be quite a challenge.
- On a TLB miss, two loads from memory will be required to get the right translation information from the page table
 - one for the page directory, and one for the PTE itself
 - time-space trade-off
- Complex

64-bit Logical Address Space

- Even two-level paging scheme not sufficient
- If page size is 4 KB (2¹²)
 - Then page table has 2⁵² entries
 - If two level scheme, inner page tables could be 2¹⁰ 4-byte entries, and outer page table has 2⁴² entries or 2⁴⁴ bytes

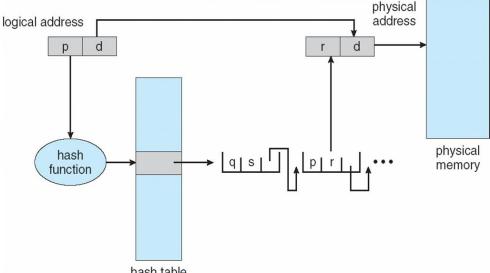
outer page	inner page	offset
p_1	p_2	d
42	10	12

- One solution is to add a 2nd outer page table
- But 2nd outer page table is still 2³⁴ bytes in size
 - And possibly 4 memory access to get to one physical memory location

2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12

Hashed Page Tables

- Common approach for handling address spaces larger than 32 bits
- Logical page number is hashed into a page table
 - This page table contains a chain of elements hashing to the same location
- Each element contains
 - (1) the logical page number
 - (2) the value of the mapped page frame
 - (3) a pointer to the next element
- Logical page numbers are compared in this chain searching for a match
 - If a match is found, the corresponding physical frame is extracted

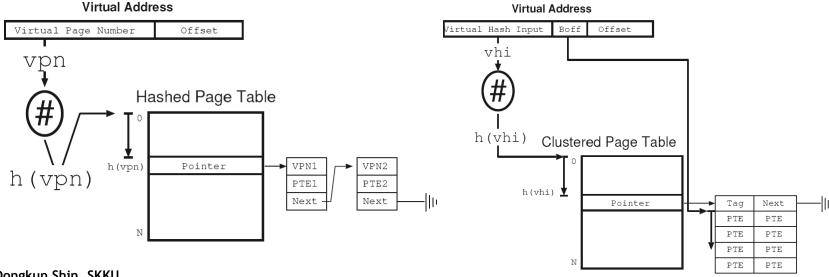


Dongkun Shin, SKKU hash table

Hashed Page Table

Clustered page table

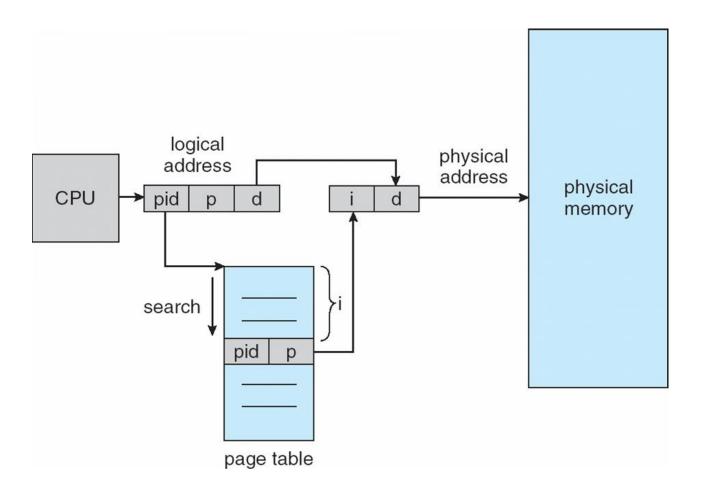
- Variation for 64-bit addresses.
- Similar to hashed but each entry refers to several pages (such as 16) rather than 1
 - Single page table entry can store the mappings for several physical page frames
- Especially useful for sparse address spaces (where memory) references are non-contiguous and scattered)



Inverted Page Table

- PowerPC
- One entry for each real page of memory
 - Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
 - Decreases memory needed to store each page table,
 - Increases time needed to search the table when a page reference occurs
 - Use hash table to limit the search to one or at most a few page-table entries
 - TLB can accelerate access
- But how to implement shared memory?
 - One mapping of a virtual address to the shared physical address

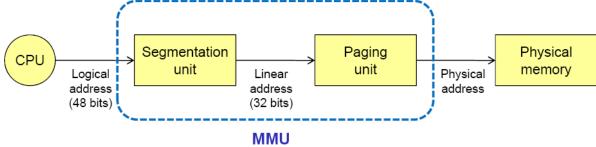
Inverted Page Table Architecture



Intel IA-32 Architecture

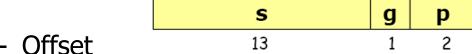
Supports both segmentation and segmentation with paging

- Each segment can be 4 GB
- Up to 16K (2¹⁴) segments per process
- Divided into two partitions
 - 1st partition: 8K (2¹³) private segments
 - kept in local descriptor table (LDT)
 - 2nd partition: 8K (2¹³) shared segments
 - kept in global descriptor table (GDT)
- Each entry of LDT and GDT
 - 8-byte segment descriptor
 - Detailed information on a particular segment
 - Base location, limit of the segment, etc

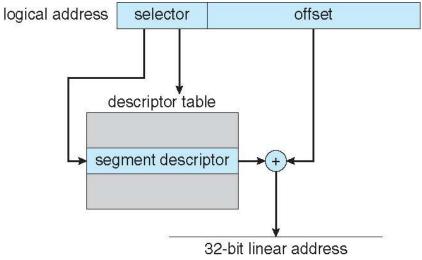


Intel IA-32 Architecture

- **CPU** generates 48 bit logical address (selector, offset)
 - Selector (s, g, p): 16 bits
 - s: segment number (13 bits)
 - g: whether the segment is in the LDT or GDT (1 bit)
 - p: deals with protection (2 bits)

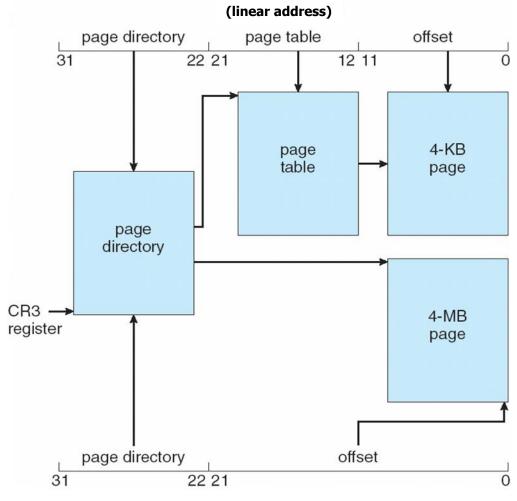


32-bit number (location within the segment)



Intel IA-32 Paging Architecture

- Page size: 4KB or 4MB
- CR3 points to the page directory for the current process



Intel x86-64

- Current generation Intel x86 architecture
- 64 bits is ginormous (> 16 exabytes)
- In practice only implement 48 bit addressing
 - Page sizes of 4 KB, 2 MB, 1 GB
 - Four levels of paging hierarchy
- Can also use PAE (Page Address Extension) so virtual addresses are 48 bits and physical addresses are 52 bits

			page map)	page directory		page		page			
١	unused	d	level 4		pointer table		directory		table		offset	
	63	48 4	7	39	38 30	29	2	1 20	,	12 11		0

ARM Architecture

- 4 KB and 16 KB pages
- 1 MB and 16 MB pages (termed sections)
- One-level paging for sections, two-level for smaller pages

